

Sabhya Chhabria

📍 San Francisco, CA | 📞 +1(650)861-0409 | ✉️ sabhyachhabria@gmail.com | [in LinkedIn](#) | [GitHub](#)

Education

M.S. Computer Science, Princeton University, GPA: 3.85
B.A. Computer Science (*Honors*), Cornell University, GPA: 3.82

Aug 2022 - May 2024
Aug 2018 - Dec 2021

Work Experience

Databricks — AI Software Engineer IV (*San Francisco, CA*)

July 2024 - Present

- Core contributor to the 0-to-1 “war room” build of [Omnigent](#) 🌀, an open-source (Apache 2.0) meta-harness that runs Claude Code, Codex, and Cursor behind a single command and unified interface; built core Client/Server/Runner runtime and native harness support across coding agents.
- Built multi-agent harness orchestration and real-time multiplayer collaboration (Polly) for parallel multi-AI coding, plus cloud sandbox execution for remote coding agents, with persistent sessions surfaced across web, mobile, and terminal.
- Led 0-to-1 development of Supervisor Agent and Knowledge Assistant, architecting FastAPI services and building agent orchestration on top of OpenAI and Anthropic Agents SDKs; delivered production-ready solution in one month, demoed at the [Data & AI Summit Keynote](#).
- Drove rapid product adoption, scaling from 0 to 1,500+ weekly active customers and managing multiple enterprise proof-of-concept engagements with Fortune 500 companies.
- Built background (asynchronous) execution for the [Supervisor API](#), enabling long-running, multi-tool agent workflows to run via a submit-and-poll model with per-call MCP tool-approval gating.
- Built comprehensive evaluation pipelines via daily CI/CD jobs to benchmark agent products. Hill-climbed Supervisor Agent quality on the OfficeQA benchmark from 21.4% to 50.2% through context engineering and agent harness development.
- Designed flexible agent-tool interface enabling customers to integrate custom Model-Context-Protocol (MCP) tools into Supervisor Agent and Knowledge Assistant products, used by over 100 customers now.
- Optimized Supervisor Agent performance, reducing time-to-first-token from $\sim 10s$ to $< 1s$ and end-to-end latency by $\sim 50\%$ through parallel tool calling and dynamic error handling, building a fault-tolerant agentic harness.

Princeton NLP — Graduate Research Assistant (*Princeton, NJ*)

Aug 2022 - May 2024

- Explored strategies to model thought trajectories of LLMs (ChatGPT & LLAMA-2) and then induce them with human-like thinking patterns through meticulous prompt optimization designed to address complex insight problems.
- Spring 2023 — Leveraged retrieval-based inference (kNN Shapley) as a surrogate for estimating Shapley scores, extending data valuation for NLP tasks and reducing computational and memory footprint by $> 90\times$ over previous methods.
- Fall 2022 — Developed a large-scale, task-agnostic datastore with an efficient encoder, optimizing data storage and retrieval for GPT-2 on downstream tasks (SuperGLUE), outperforming previous state of the art results by an average of 2%.

Scale AI — Machine Learning Research Intern (*San Francisco, CA*)

May 2022 - Aug 2022

- Built a large-scale document-extraction pipeline for fixed-format documents (W2s, I9s) and established LayoutLM baselines (0.85 F1) on Named Entity Recognition; reframed NER as a 2D object detection task and finetuned YOLOv5 to a competitive 0.82 F1 through hyperparameter tuning.

Snap Inc — Software Engineer Intern (*Santa Monica, CA*)

May 2021 - Aug 2021

- Engineered a concurrent model-assessment pipeline running 1000+ experiments and pioneered ad-ranking reliability metrics (L1/L2 distance, SMAPE, Jensen-Shannon divergence) to assess the reproducibility of production models.

Skills

Languages	Python, Go, TypeScript, Java, C/C++, SQL, Bash
Agentic Systems	MCP, Claude Code, Codex, OpenAI & Anthropic Agents SDKs, multi-agent orchestration, agent harness design, tool calling, evals & benchmarking
ML/AI	PyTorch, Hugging Face Transformers, MLflow, LangChain, CUDA
Infrastructure	Docker, Kubernetes, AWS, GCP, Spark, Databricks, FastAPI, Modal, Daytona, Git, CI/CD